



main ▾

[examples](#) / [cpp](#) / [autograd](#) / [autograd.cpp](#) 

Expand file tree

Will Feng C++ autograd example (#745)

ad775ac · 4 years ago

History

**Code**

Blame 191 lines (147 loc) · 5.86 KB

Raw



```
1  #include <torch/torch.h>
2  #include <iostream>
3
4  using namespace torch::autograd;
5
6  void basic_autograd_operations_example() {
7      std::cout << "==== Running: \"Basic autograd operations\" =====> << std::endl;
8
9      // Create a tensor and set ``torch::requires_grad()`` to track computation with it
10     auto x = torch::ones({2, 2}, torch::requires_grad());
11     std::cout << x << std::endl;
12
13     // Do a tensor operation:
14     auto y = x + 2;
15     std::cout << y << std::endl;
16
17     // ``y`` was created as a result of an operation, so it has a ``grad_fn``.
18     std::cout << y.grad_fn()->name() << std::endl;
19
20     // Do more operations on ``y``
21     auto z = y * y * 3;
```

```

22     auto out = z.mean();
23
24     std::cout << z << std::endl;
25     std::cout << z.grad_fn()->name() << std::endl;
26     std::cout << out << std::endl;
27     std::cout << out.grad_fn()->name() << std::endl;
28
29     // ``.requires_grad_( ... )`` changes an existing tensor's ``requires_grad`` flag in-place.
30     auto a = torch::randn({2, 2});
31     a = ((a * 3) / (a - 1));
32     std::cout << a.requires_grad() << std::endl;
33
34     a.requires_grad_(true);
35     std::cout << a.requires_grad() << std::endl;
36
37     auto b = (a * a).sum();
38     std::cout << b.grad_fn()->name() << std::endl;
39
40     // Let's backprop now. Because ``out`` contains a single scalar, ``out.backward()``
41     // is equivalent to ``out.backward(torch::tensor(1.))``.
42     out.backward();
43
44     // Print gradients d(out)/dx
45     std::cout << x.grad() << std::endl;
46
47     // Now let's take a look at an example of vector-Jacobian product:
48     x = torch::randn(3, torch::requires_grad());
49
50     y = x * 2;
51     while (y.norm().item<double>() < 1000) {
52         y = y * 2;
53     }
54
55     std::cout << y << std::endl;
56     std::cout << y.grad_fn()->name() << std::endl;
57
58     // ...

```

```

58 // If we want the vector-Jacobian product, pass the vector to ``backward`` as argument:
59 auto v = torch::tensor({0.1, 1.0, 0.0001}, torch::kFloat);
60 y.backward(v);
61
62 std::cout << x.grad() << std::endl;
63
64 // You can also stop autograd from tracking history on tensors that require gradients
65 // either by putting ``torch::NoGradGuard`` in a code block
66 std::cout << x.requires_grad() << std::endl;
67 std::cout << x.pow(2).requires_grad() << std::endl;
68
69 {
70     torch::NoGradGuard no_grad;
71     std::cout << x.pow(2).requires_grad() << std::endl;
72 }
73
74 // Or by using ``.detach()`` to get a new tensor with the same content but that does
75 // not require gradients:
76 std::cout << x.requires_grad() << std::endl;
77 y = x.detach();
78 std::cout << y.requires_grad() << std::endl;
79 std::cout << x.eq(y).all().item<bool>() << std::endl;
80 }
81
82 void compute_higher_order_gradients_example() {
83     std::cout << "==== Running \"Computing higher-order gradients in C++\" =====> << std::endl;
84
85     // One of the applications of higher-order gradients is calculating gradient penalty.
86     // Let's see an example of it using ``torch::autograd::grad``:
87
88     auto model = torch::nn::Linear(4, 3);
89
90     auto input = torch::randn({3, 4}).requires_grad_(true);
91     auto output = model(input);
92
93     // Calculate loss

```

```

94     auto target = torch::randn({3, 3});
95     auto loss = torch::nn::MSELoss()(output, target);
96
97     // Use norm of gradients as penalty
98     auto grad_output = torch::ones_like(output);
99     auto gradient = torch::autograd::grad({output}, {input}, /*grad_outputs=*/{grad_output}, /*create_graph=*/true)[0];
100    auto gradient_penalty = torch::pow((gradient.norm(2, /*dim=*/1) - 1), 2).mean();
101
102    // Add gradient penalty to loss
103    auto combined_loss = loss + gradient_penalty;
104    combined_loss.backward();
105
106    std::cout << input.grad() << std::endl;
107 }
108
109 // Inherit from Function
110 class LinearFunction : public Function<LinearFunction> {
111 public:
112     // Note that both forward and backward are static functions
113
114     // bias is an optional argument
115     static torch::Tensor forward(
116         AutogradContext *ctx, torch::Tensor input, torch::Tensor weight, torch::Tensor bias = torch::Tensor()) {
117         ctx->save_for_backward({input, weight, bias});
118         auto output = input.mm(weight.t());
119         if (bias.defined()) {
120             output += bias.unsqueeze(0).expand_as(output);
121         }
122         return output;
123     }
124
125     static tensor_list backward(AutogradContext *ctx, tensor_list grad_outputs) {
126         auto saved = ctx->get_saved_variables();
127         auto input = saved[0];
128         auto weight = saved[1];
129         auto bias = saved[2];
130     }

```

```

130
131     auto grad_output = grad_outputs[0];
132     auto grad_input = grad_output.mm(weight);
133     auto grad_weight = grad_output.t().mm(input);
134     auto grad_bias = torch::Tensor();
135     if (bias.defined()) {
136         grad_bias = grad_output.sum(0);
137     }
138
139     return {grad_input, grad_weight, grad_bias};
140 }
141 };
142
143 ✓ class MulConstant : public Function<MulConstant> {
144     public:
145     ✓ static torch::Tensor forward(AutogradContext *ctx, torch::Tensor tensor, double constant) {
146         // ctx is a context object that can be used to stash information
147         // for backward computation
148         ctx->saved_data["constant"] = constant;
149         return tensor * constant;
150     }
151
152     ✓ static tensor_list backward(AutogradContext *ctx, tensor_list grad_outputs) {
153         // We return as many input gradients as there were arguments.
154         // Gradients of non-tensor arguments to forward must be `torch::Tensor()`.
155         return {grad_outputs[0] * ctx->saved_data["constant"].toDouble(), torch::Tensor()};
156     }
157 };
158
159 ✓ void custom_autograd_function_example() {
160     std::cout << "==== Running \"Using custom autograd function in C++\" =====> << std::endl;
161     {
162         auto x = torch::randn({2, 3}).requires_grad_();
163         auto weight = torch::randn({4, 3}).requires_grad_();
164         auto y = LinearFunction::apply(x, weight);
165         y.sum().backward();
166     }
167 }

```

```
166
167     std::cout << x.grad() << std::endl;
168     std::cout << weight.grad() << std::endl;
169 }
170 {
171     auto x = torch::randn({2}).requires_grad_();
172     auto y = MulConstant::apply(x, 5.5);
173     y.sum().backward();
174
175     std::cout << x.grad() << std::endl;
176 }
177 }
178
179 ✓ int main() {
180     std::cout << std::boolalpha;
181
182     basic_autograd_operations_example();
183
184     std::cout << "\n";
185
186     compute_higher_order_gradients_example();
187
188     std::cout << "\n";
189
190     custom_autograd_function_example();
191 }
```